

# Distributed Data Processing

UA.DETI.CBD  
José Luis Oliveira / Carlos Costa



1

## Introduction

- ❖ Large-Scale Data Processing
  - aim to use 1000s of CPUs, but with simplified managing
- ❖ MapReduce
- ❖ Apache Hadoop
- ❖ To start:
  - <https://www.youtube.com/watch?v=9s-vSeWej1U>



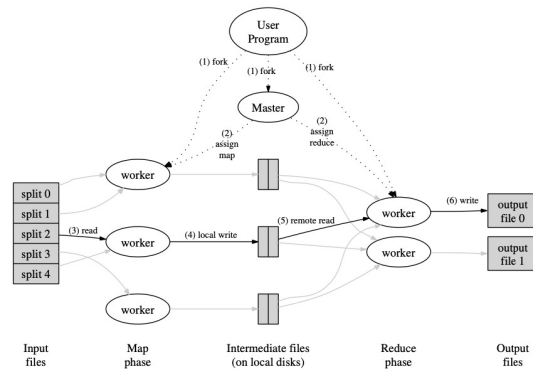
2

2

## What is MapReduce?

- ❖ "MapReduce is a programming model and an associated implementation for processing and generating large data sets"

- Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: Simplified data processing on large clusters." (2004). [Google]



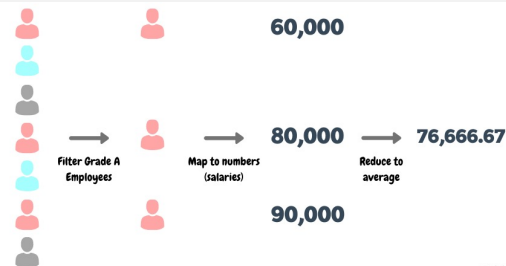
## What is MapReduce?

- ❖ Terms are borrowed from Functional Language (e.g., Lisp)
- ❖ Example: Sum of squares
  - **map** square '(1 2 3 4))
    - Output: (1 4 9 16)
    - processes each record sequentially and independently
  - **reduce** + '(1 4 9 16))
    - (+ 16 (+ 9 (+ 4 1)))
    - Output: 30
    - processes set of all records in batches
- ❖ This concept has been reused in several programming languages and computational tools

## Examples: Java stream

```
List<String> values = Arrays.asList("1","2","3","4","5","6","7");
int soma = values.stream()
    .map(num->Integer.parseInt(num))
    .reduce(0, Integer::sum);
```

```
OptionalDouble avgSalary = employees.stream()
    .filter(e -> e.grade == 'A') // filter 'A' grade employees
    .mapToInt((e) -> e.salary)   // get IntStream of salaries
    .average();                  // get average
```



5

5

## Examples: MongoDB

Database:

```
db.orders.insertMany([
  { _id: 1, cust_id: "Ant O. Knee", ord_date: new Date("2020-03-01"), price: 25,
    items: [{sku: "oranges", qty: 5, price: 2.5}, {sku: "apples", qty: 5, price: 2.5}],
    status: "A" },
  { _id: 2, cust_id: "Ant O. Knee", ord_date: new Date("2020-03-08"), price: 70, items: [
    {sku: "oranges", qty: 8, price: 2.5}, {sku: "chocolates", qty: 5, price: 10}], status:
    "A" },
  ...
])
```

Return the Total Price Per Customer:

```
var mapFunction1 = function() {
  emit(this.cust_id, this.price);
};
var reduceFunction1 = function(keyCustId, valuesPrices) {
  return Array.sum(valuesPrices);
};
db.orders.mapReduce(
  mapFunction1,
  reduceFunction1,
  { out: "map_reduce_example" }
)
db.map_reduce_example.find().sort({ _id: 1 })
```

```
{ "_id": "Ant O. Knee", "value": 95 }
{ "_id": "Busby Bee", "value": 125 }
{ "_id": "Cam Elot", "value": 60 }
...
```



<https://docs.mongodb.com/manual/tutorial/map-reduce-examples/>

6

6

## MapReduce frameworks

- ❖ Library/Tools that allow easily writing applications that process large amounts of data in parallel
  - distributed across several nodes
- ❖ Good retry/failure semantics
- ❖ Solution Pattern
  - many problems can be modelled in this way
- ❖ Implementations
  - **Hadoop**: the mapper and reducer are each a Java class that implements a particular interface.
  - **Spark**: newer and faster, it processes data in RAM using a concept known as an RDD, Resilient Distributed Dataset.



7

7

## MapReduce frameworks

- ❖ **Framework...**
  - Automatic parallelization & distribution
  - Fault tolerance
  - I/O scheduling
  - Monitoring & status updates
- ❖ Two main tasks:
  - Mappers & Reducers
- ❖ Pipeline
  - Splits the input data-set into independent chunks
  - Mappers process of chunks in a parallel manner
  - aggregates the outputs of the maps
  - Reducers process the mappers output




8

8

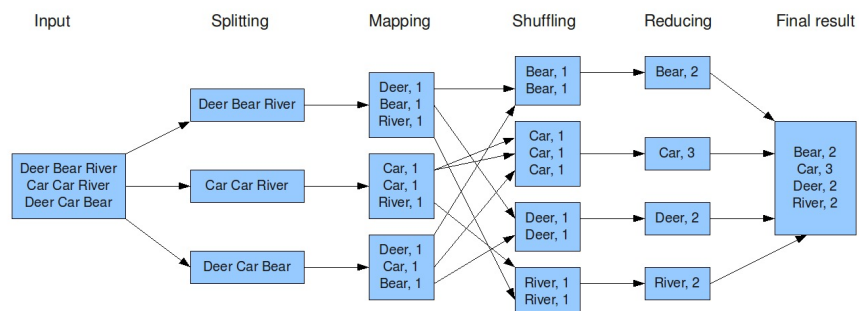
## MapReduce dataflow

- ❖ The **map** function
  - is called once for every input record to extract (one or more) key-value from the input record
- ❖ MapReduce framework
  - collects all the key-value pairs with the same key – **shuffle**
- ❖ The **reduce** function
  - iterates over each collection of values with the same key and can produce output records
    - e.g., the number of occurrences of the key



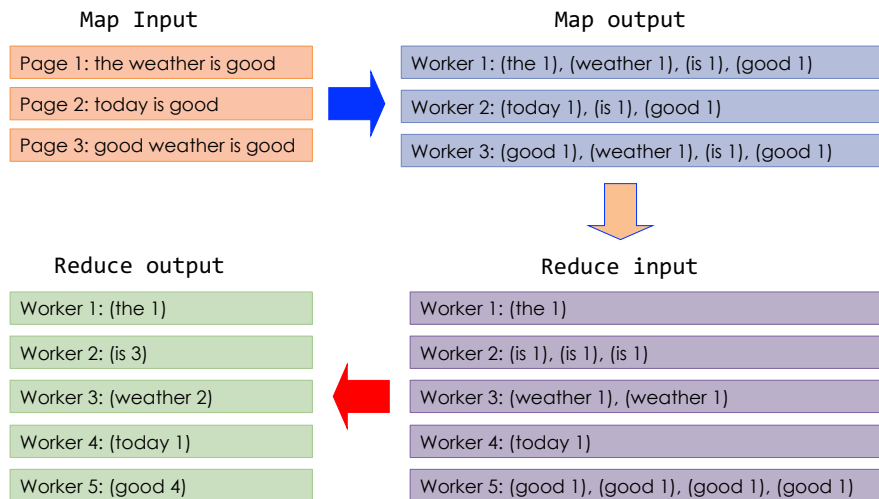
9

## Example: Words count dataflow



10

## Words count – tasks execution



11

## MapReduce applications

- ❖ words count / histogram
- ❖ distributed search
- ❖ distributed sort
- ❖ web link-graph reversal
- ❖ term-vector per host
- ❖ web access log stats
- ❖ inverted index construction
- ❖ document clustering
- ❖ machine learning
- ❖ statistical machine translation
- ❖ ...


12

# Apache Hadoop

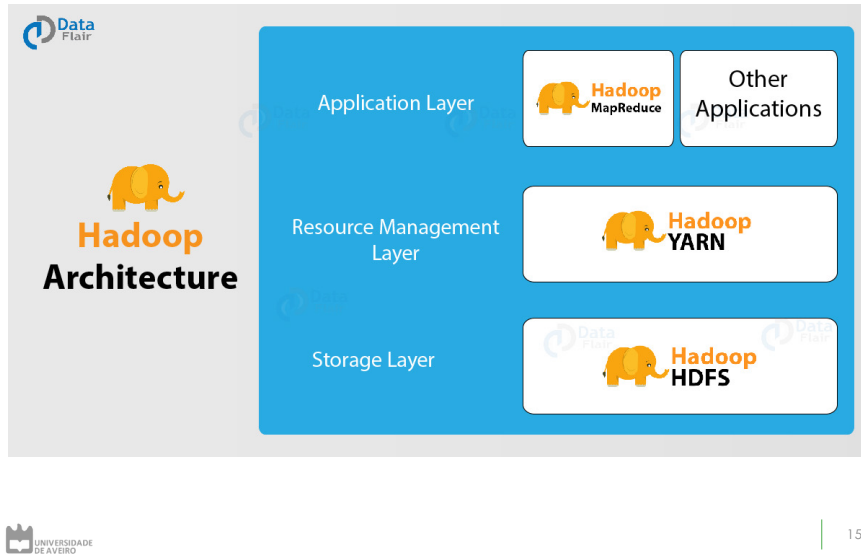


## Apache Hadoop Framework

- ❖ A framework that allows distributed processing of **large data sets across clusters of computers** using simple programming models.
- ❖ Open-source framework
  - Java
  - <https://hadoop.apache.org>
- ❖ **Main components**
  - Hadoop Distributed File System (**HDFS**)
    - Distributed, scalable, and portable file system
  - Hadoop Yet Another Resource Negotiator (**YARN**)
  - Hadoop Common Utilities
  - Hadoop **MapReduce**
    - Implementation of the MapReduce programming model



## Apache Hadoop Framework



15

## Hadoop Distributed File System (HDFS)

- ❖ Distributed file system designed to run on commodity hardware.
- ❖ Stores data redundantly on multiple nodes – Fault-tolerant file system
  - 3+ replicas for each block
  - Default Block Size : 128MB
- ❖ Master-Slave architecture
  - NameNode: Master
  - DataNode: Slave
- ❖ Typical usage pattern
  - Huge files (GB to TB)
  - Data is rarely updated
  - Reads and appends are common
    - Usually, random read/write operations are not performed

16



## HDFS – Assumptions and Goals

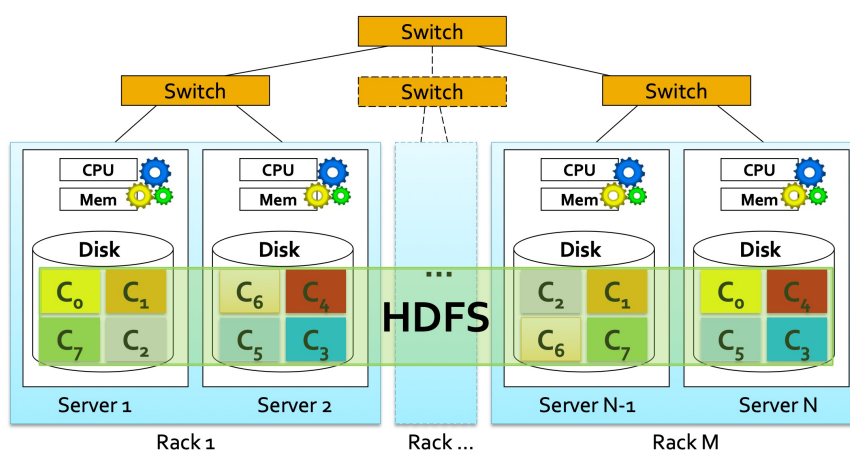
- ❖ **Hardware Failure**
  - An HDFS instance may consist of +100/+1000 of server machines.
- ❖ **Streaming Data Access**
  - HDFS is designed more for batch processing rather than interactive use. The emphasis is on high throughput of data access rather than low latency of data access.
- ❖ **Large Data Sets**
  - Applications that run on HDFS have large data sets. A typical file in HDFS is gigabytes to terabytes in size.
- ❖ **Simple Coherency Model**
  - HDFS applications need a write-once-read-many access model for files. A file once created, written, and closed need not be changed except for appends and truncates.
- ❖ **“Moving Computation is Cheaper than Moving Data”**
  - Especially true when the size of the data set is huge. Minimizes network congestion and increases the throughput of the system.
- ❖ **Portability Across Heterogeneous Platforms**



17

17

## HDFS



Example with number of replicas per chunk = 2



18

18

## HDFS

### ❖ Master-Slave architecture

- Master: NameNode
- Slave: DataNode

### ❖ The Master node is a special node/server that

- Stores HDFS metadata
  - E.g., the mapping between the name of a file and the location of its chunks
- Might be replicated

### ❖ Client applications: file access through HDFS APIs

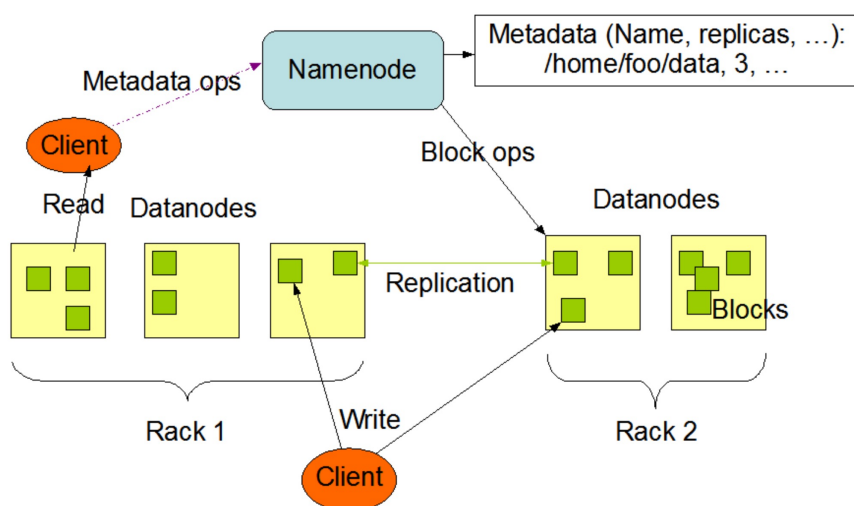
- Talk to the master node to find data/chunk servers associated with the file of interest
- Connect to the selected chunk servers to access data



19

19

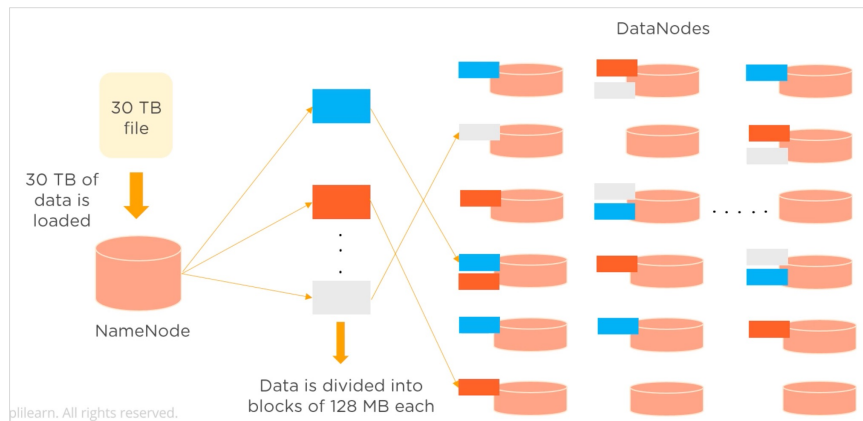
## Interaction between HDFS components



20

20

## Storage example



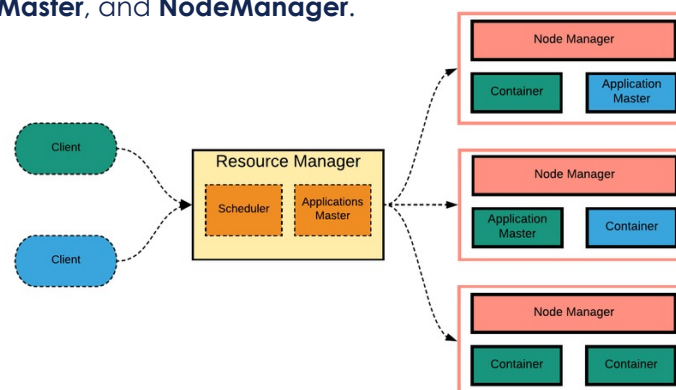
<https://www.youtube.com/watch?v=iANBytZ26MI>

21

21

## Hadoop YARN

- ❖ Resource management and job scheduling.
  - Global **ResourceManager**, per-application **Application Master**, and **NodeManager**.



22

22

## Hadoop MapReduce

- ❖ MapReduce is the data processing layer of Hadoop.
- ❖ MapReduce job comprises a number of map tasks and reduces tasks.
  - Each task works on a part of data. This distributes the load across the cluster.
  - **Map task** load, parse, transform and filter data.
  - Each **reduce task** works on the sub-set of output from the map tasks (e.g., by applying grouping and aggregation).



23

23

## Apache Hadoop

- ❖ Install & config
- ❖ Run
  - ~ `sbin ./start-dfs.sh`
- ❖ Basic help for all the commands
  - ~ `hadoop`
- ❖ Distributed file system commands
  - ~ `hadoop fs`
- ❖ Execution of MapReduce jobs
  - ~ `hadoop jar`



24

24

## HDFS commands

```

~ hadoop fs [generic options]
  [-appendToFile <localsrc> ... <dst>]
  [-cat [-ignoreCrc] <src> ...]
  [-checksum [-v] <src> ...]
  [-chgrp [-R] GROUP PATH...]
  [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
  [-chown [-R] [OWNER][:[GROUP]] PATH...]
  [-copyFromLocal [-f] [-p] [-l] [-d] [-t <thread count>] <localsrc> ... <dst>]
  [-copyToLocal [-f] [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
  [-count [-q] [-h] [-v] [-t <storage type>]] [-u] [-x] [-e] <path> ...]
  [-cp [-f] [-p | -p[topax]] [-d] <src> ... <dst>]
  [-createSnapshot <snapshotDir> [<snapshotName>]]
  [-deleteSnapshot <snapshotDir> <snapshotName>]
  [-df [-h] [<path> ...]]
  [-du [-s] [-h] [-v] [-x] <path> ...]
  [-expunge [-immediate] [-fs <path>]]
  [-find <path> ... <expression> ...]
  [-get [-f] [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
  ...

```



25

25

## Hadoop – HDFS Commands

- ❖ Help
  - ~ `hadoop fs -help`
- ❖ Creating a directory
  - ~ `hadoop fs -mkdir /myhdfsdir`
- ❖ Listing a directory
  - ~ `hadoop fs -ls`
  - ~ `hadoop fs -ls /myhdfsdir`
- ❖ Copy a file/dir from local file system to HDFS
  - ^ `hadoop fs -put /myhdfsdir/ /user/xpto/`
- ❖ Get a file/dir from HDFS to local file system
  - ~ `hadoop fs -get /user/xpto/ /myhdfsdir/`
  - ... and many more

[https://www.tutorialspoint.com/hadoop/hadoop\\_command\\_reference.htm](https://www.tutorialspoint.com/hadoop/hadoop_command_reference.htm)



26

26

## Execute a MapReduce job

### ❖ Run the job

```
~ hadoop jar WordCount.jar /myhdfsdir/myapp/input1 /myhdfsdir/myapp/output1
```

### ❖ Retrieve and explore the job result

```
~ hadoop fs -copyToLocal /myhdfsdir/myapp/output1/part-r-00000 result.txt  
~ cat result.txt
```

### ❖ Clean the output HDFS directory

```
~ hadoop fs -rm /myhdfsdir/myapp/output1/
```

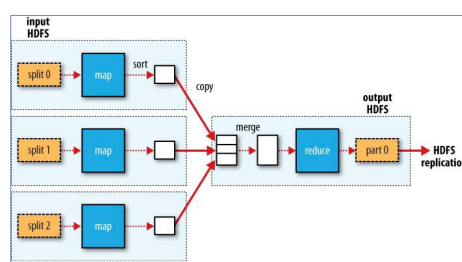


27


27

## Task execution - distinct configurations

single reduce task



multiple reduce task



29

29

## Java Interface

### ❖ Mapper class

- Implementation of the **map function**
- Template parameters
  - KEYIN, VALUEIN** – types of input key-value pairs
  - KEYOUT, VALUEOUT** – types of intermediate key-value pairs
- Intermediate pairs are emitted via **context.write(k, v)**

```
class MyMapper extends Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {
    @Override
    public void map(KEYIN key, VALUEIN value, Context context)
        throws IOException, InterruptedException
    {
        // Implementation
    }
}
```



32

32

## Java Interface

### ❖ Reducer class

- Implementation of the **reduce function**
- Template parameters
  - KEYIN, VALUEIN** – types of intermediate key-value pairs
  - KEYOUT, VALUEOUT** – types of output key-value pairs
- Output pairs are emitted via **context.write(k, v)**

```
class MyReducer extends Reducer<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {
    @Override
    public void reduce(KEYIN key, Iterable<VALUEIN> values, Context context)
        throws IOException, InterruptedException
    {
        // Implementation
    }
}
```



33

33

## MapReduce – WordCount example (1)

```
public class WordCount extends Configured implements Tool {
    private static IntWritable ONE = new IntWritable(1);
    @Override
    public int run(String[] args) throws Exception {
        FileSystem fs = FileSystem.get(getConf());
        Job job = Job.getInstance(getConf());
        job.setJarByClass(WordCount.class);
        job.setJobName("WordCount");
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setMapperClass(Map.class); ← Mapper class
        job.setCombinerClass(Reducer.class); ← Reducer class
        job.setReducerClass(Reducer.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.setInputPaths(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        if (fs.exists(new Path(args[1])))
            fs.delete(new Path(args[1]), true);
        boolean success = job.waitForCompletion(true);
        return success ? 0 : 1;
    }
}
```

Configurations



34

34

## MapReduce – WordCount example (2)

```
...
public static void main(String[] args) throws Exception {
    int ret = ToolRunner.run(new WordCount(), args);
    System.exit(ret);
}

public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException {
        String line = value.toString();
        String[] words = line.split("\\s+");
        for (String word : words) {
            if (word.trim().length() > 0) {
                Text text = new Text();
                text.set(word);
                context.write(text, ONE);
            }
        }
    }
} // end class Map
```

Mapper



35

35



## MapReduce – WordCount example (3)

```

...
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values)
            sum += val.get();
        context.write(key, new IntWritable(sum));
    }
} // end class Reduce
} // end class WordCount

```

Reducer



36

36

## MapReduce – Run the example

### Compile

```
> $ javac -classpath $CLASSPATH -d WordCountDir WordCount.java
```

### Create JAR

```
> $ jar -cvf WordCount.jar -C WordCountDir/ .
```

### Run WordCount App

```
> $ hadoop jar WordCount.jar /in/test.txt /out/wc
```

### See the Result

```
> $ hadoop fs -cat /out/wc/part-r-00000
```

```

ABOUT 1
ACCOUNT 2
ACTUAL 1
ADDITIONAL 1
ADVANCING 2
...

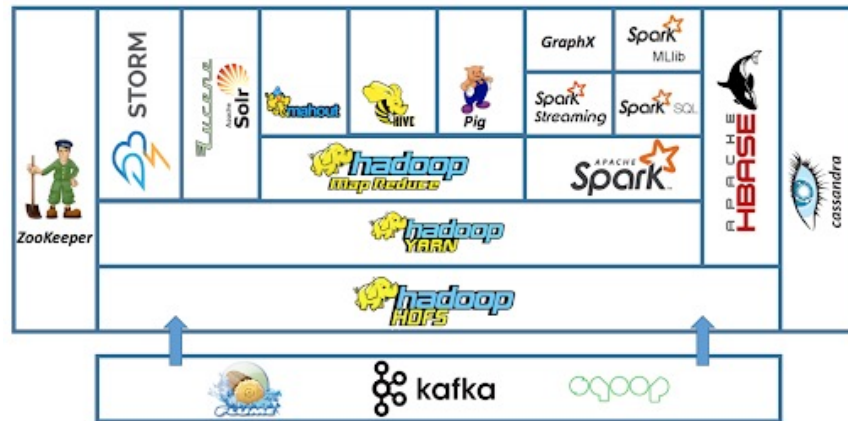
```



37

37

## Hadoop ecosystem



38

## Summary

- ❖ Concept
- ❖ MapReduce framework
  - Programming Model
  - Execution Plan
  - Dataflow
- ❖ MapReduce Applications
- ❖ Apache Hadoop
  - HDFS
  - MapReduce Programming

39